

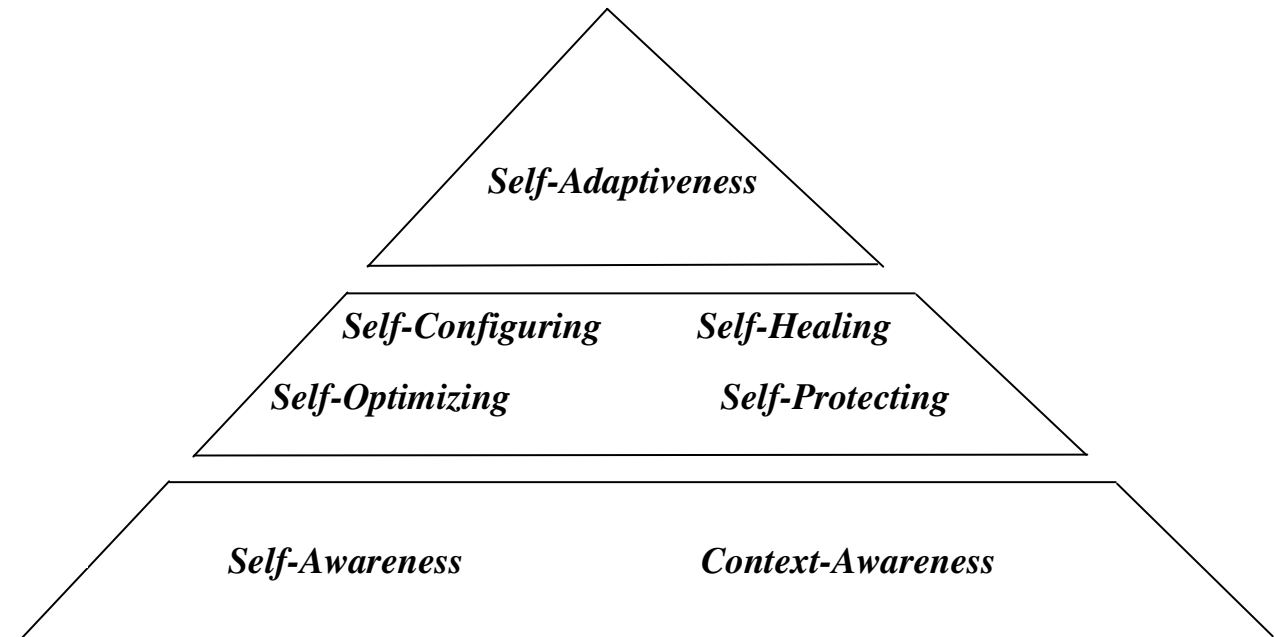

Online Learning for Self-Healing and Self-Optimization

Project Seminar

Prof. Dr. Holger Giese, Christian Medeiros Adriano , Sona Ghahremani

System Analysis & Modeling Group,
Hasso Plattner Institute for Digital Engineering
University of Potsdam, Germany

holger.giese@hpi.de
christian.adriano@hpi.de
sona.ghahremani@hpi.de



Self-healing: is the capability of discovering, diagnosing, and reacting to disruptions at runtime.

Self-optimizing: is the capability of managing performance and resource allocation in order to satisfy the requirements.

Any component can (1) crash, (2) throw exceptions, or (3) it can be destroyed (severe crash).

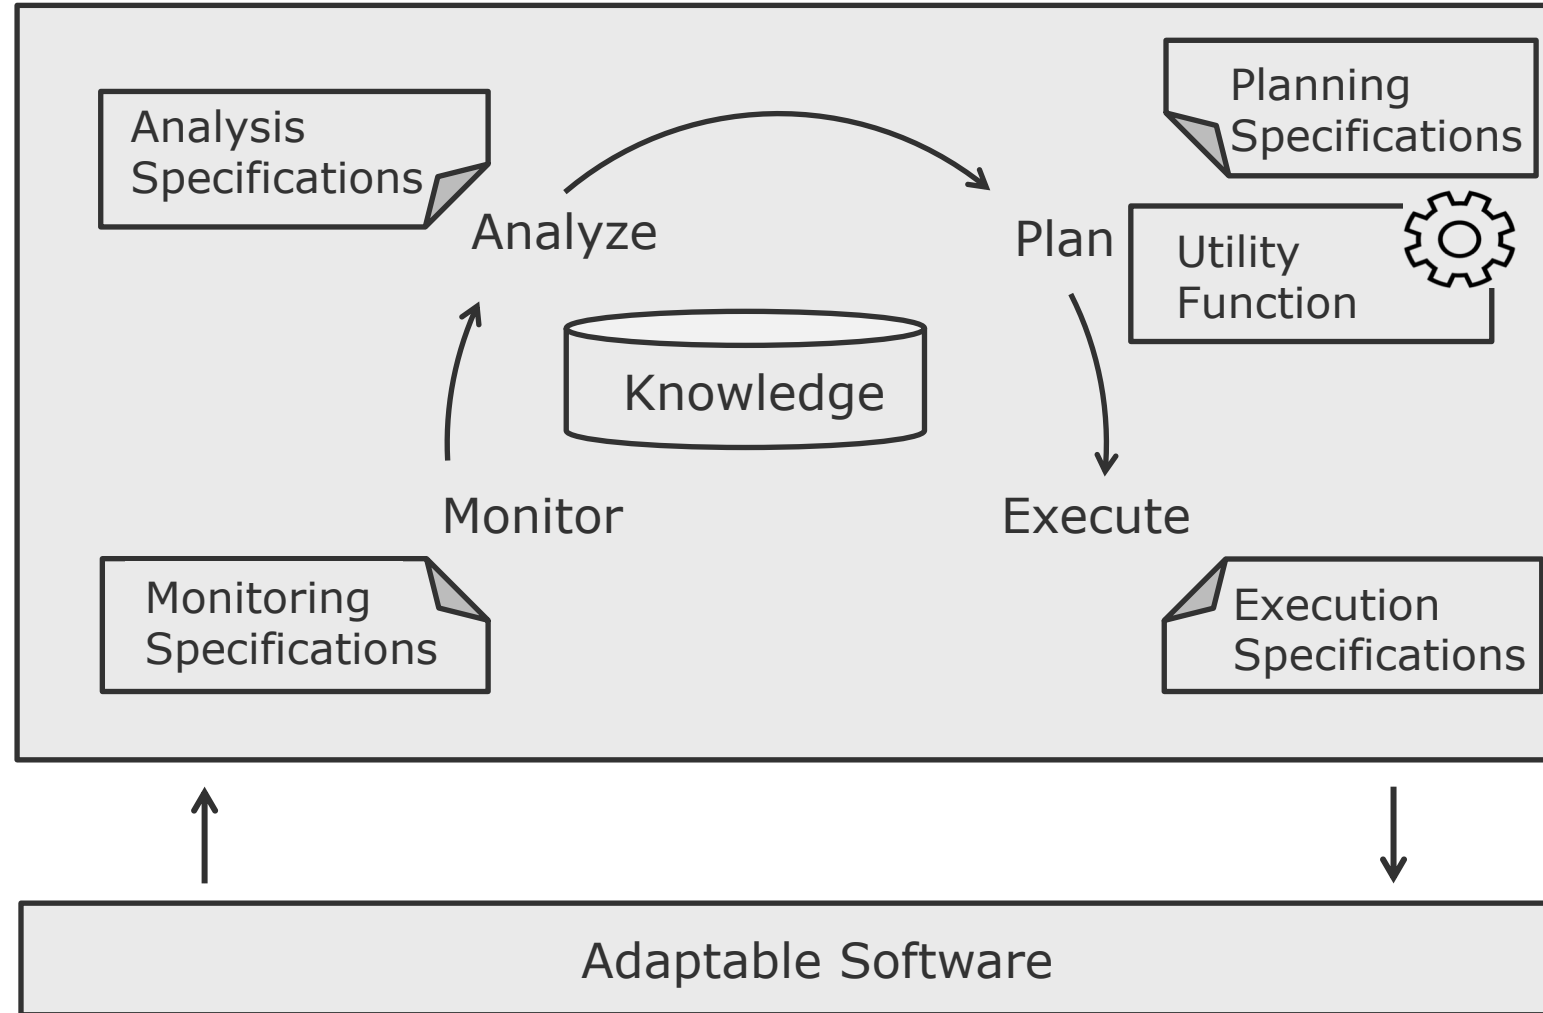
Possible countermeasures (possible actions or rules):


- light-weight redeployment of the component,
- restart the component,
- heavy-weight redeployment of the component, and
- replace a component with an alternative one, for instance, switch from internal authorization to an external authorization such as Facebook.

The load of the system can vary and result in an increased or decreased demand for each service and its replicas.

Possible countermeasures (possible actions or rules):

- Add replica to services in high demand
- Remove replica to reduce cost.






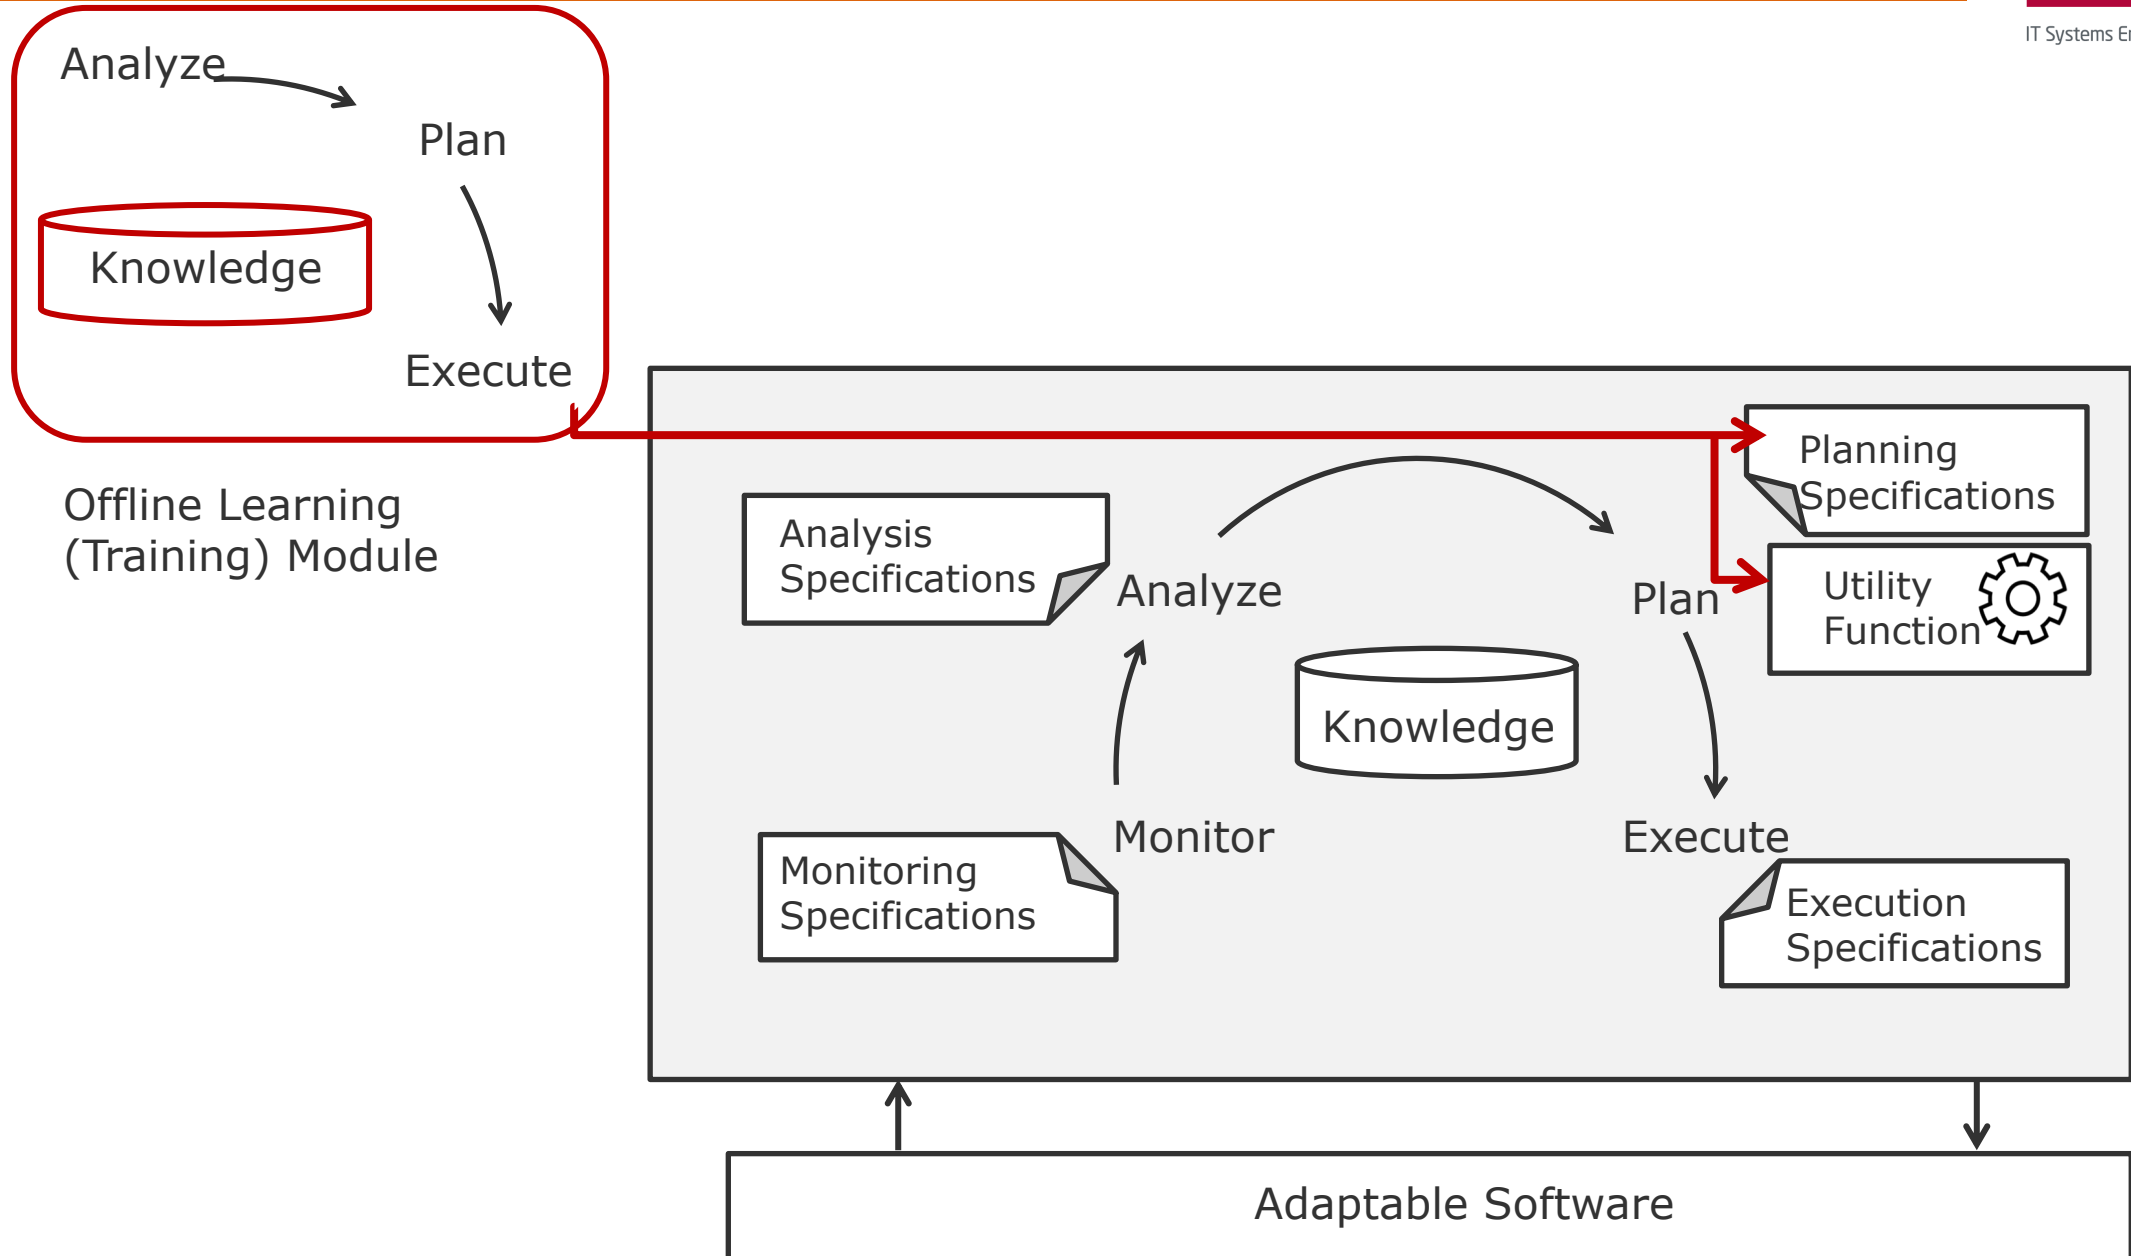
Sequence of repair/optimization steps -> **failure free / optimal solution space**

Final configuration with **highest utility**

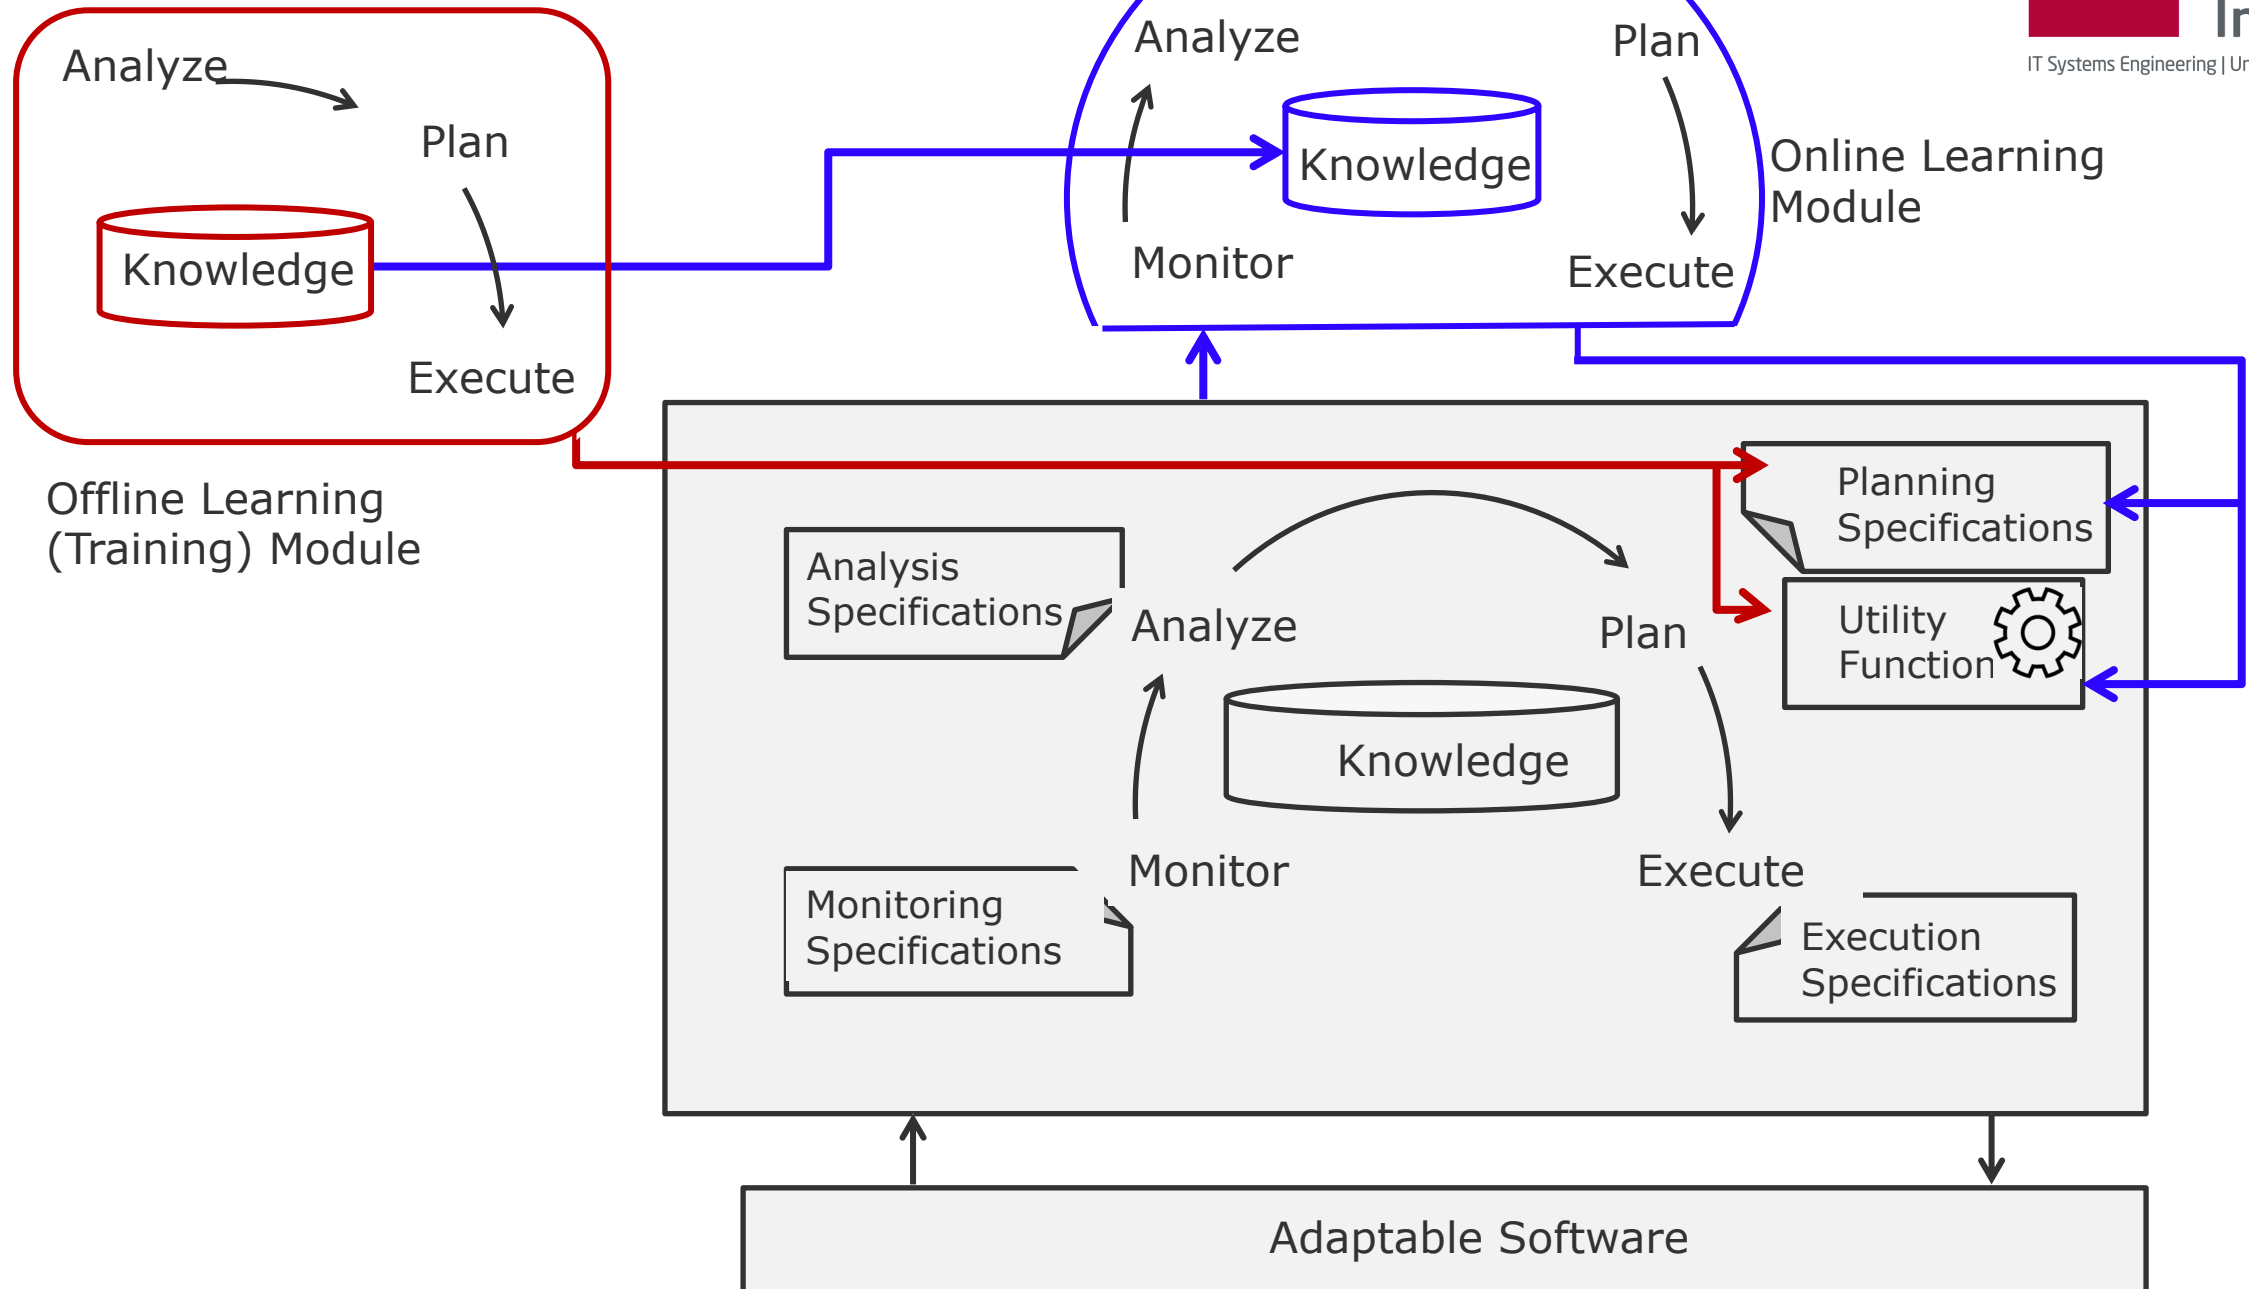
Optimal order of rules-> highest Reward



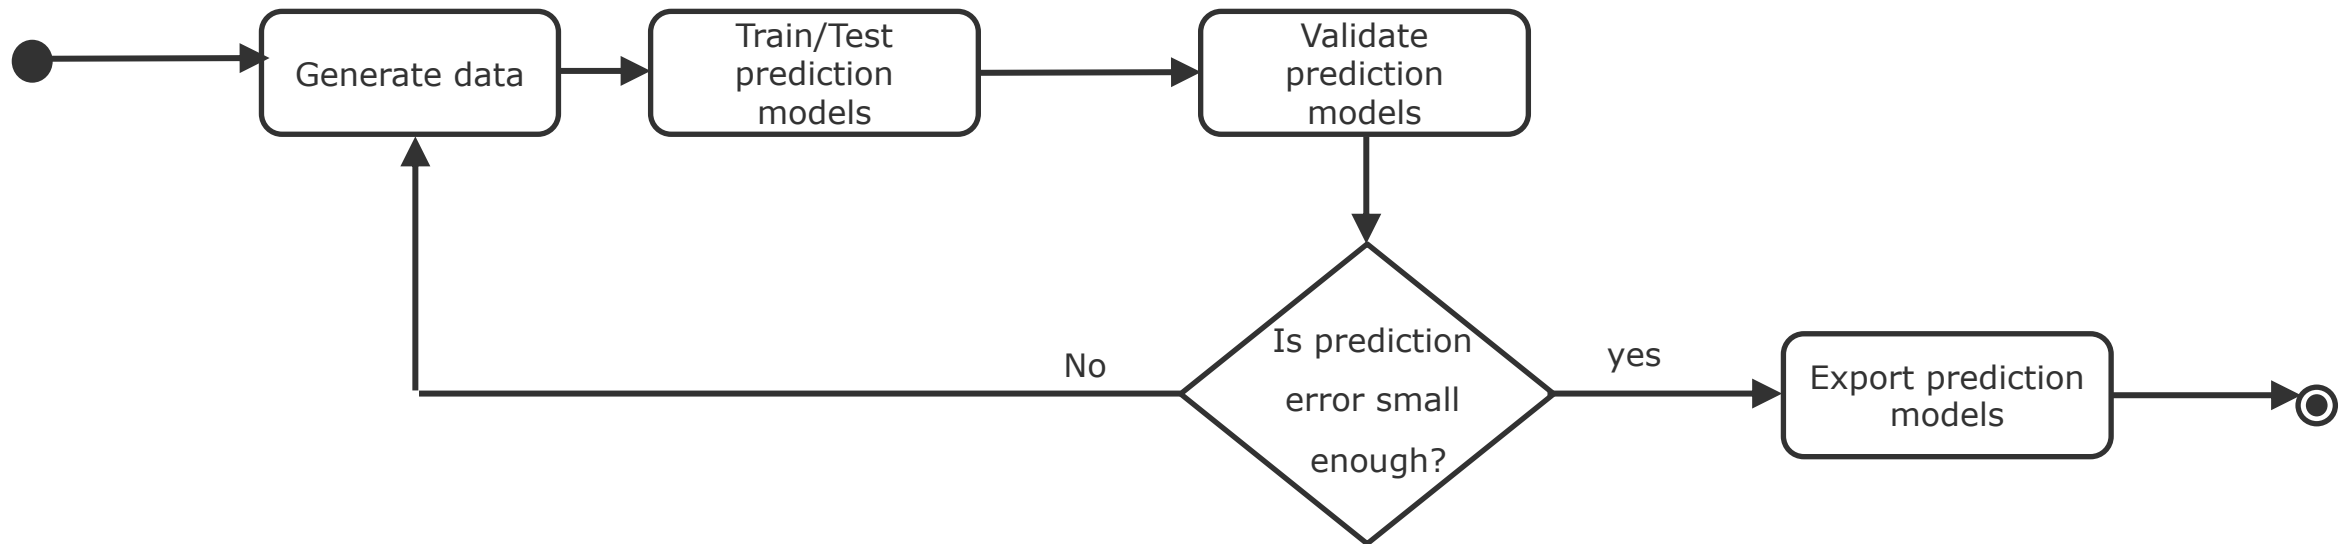
Offline Learning for Self-healing and self-optimization



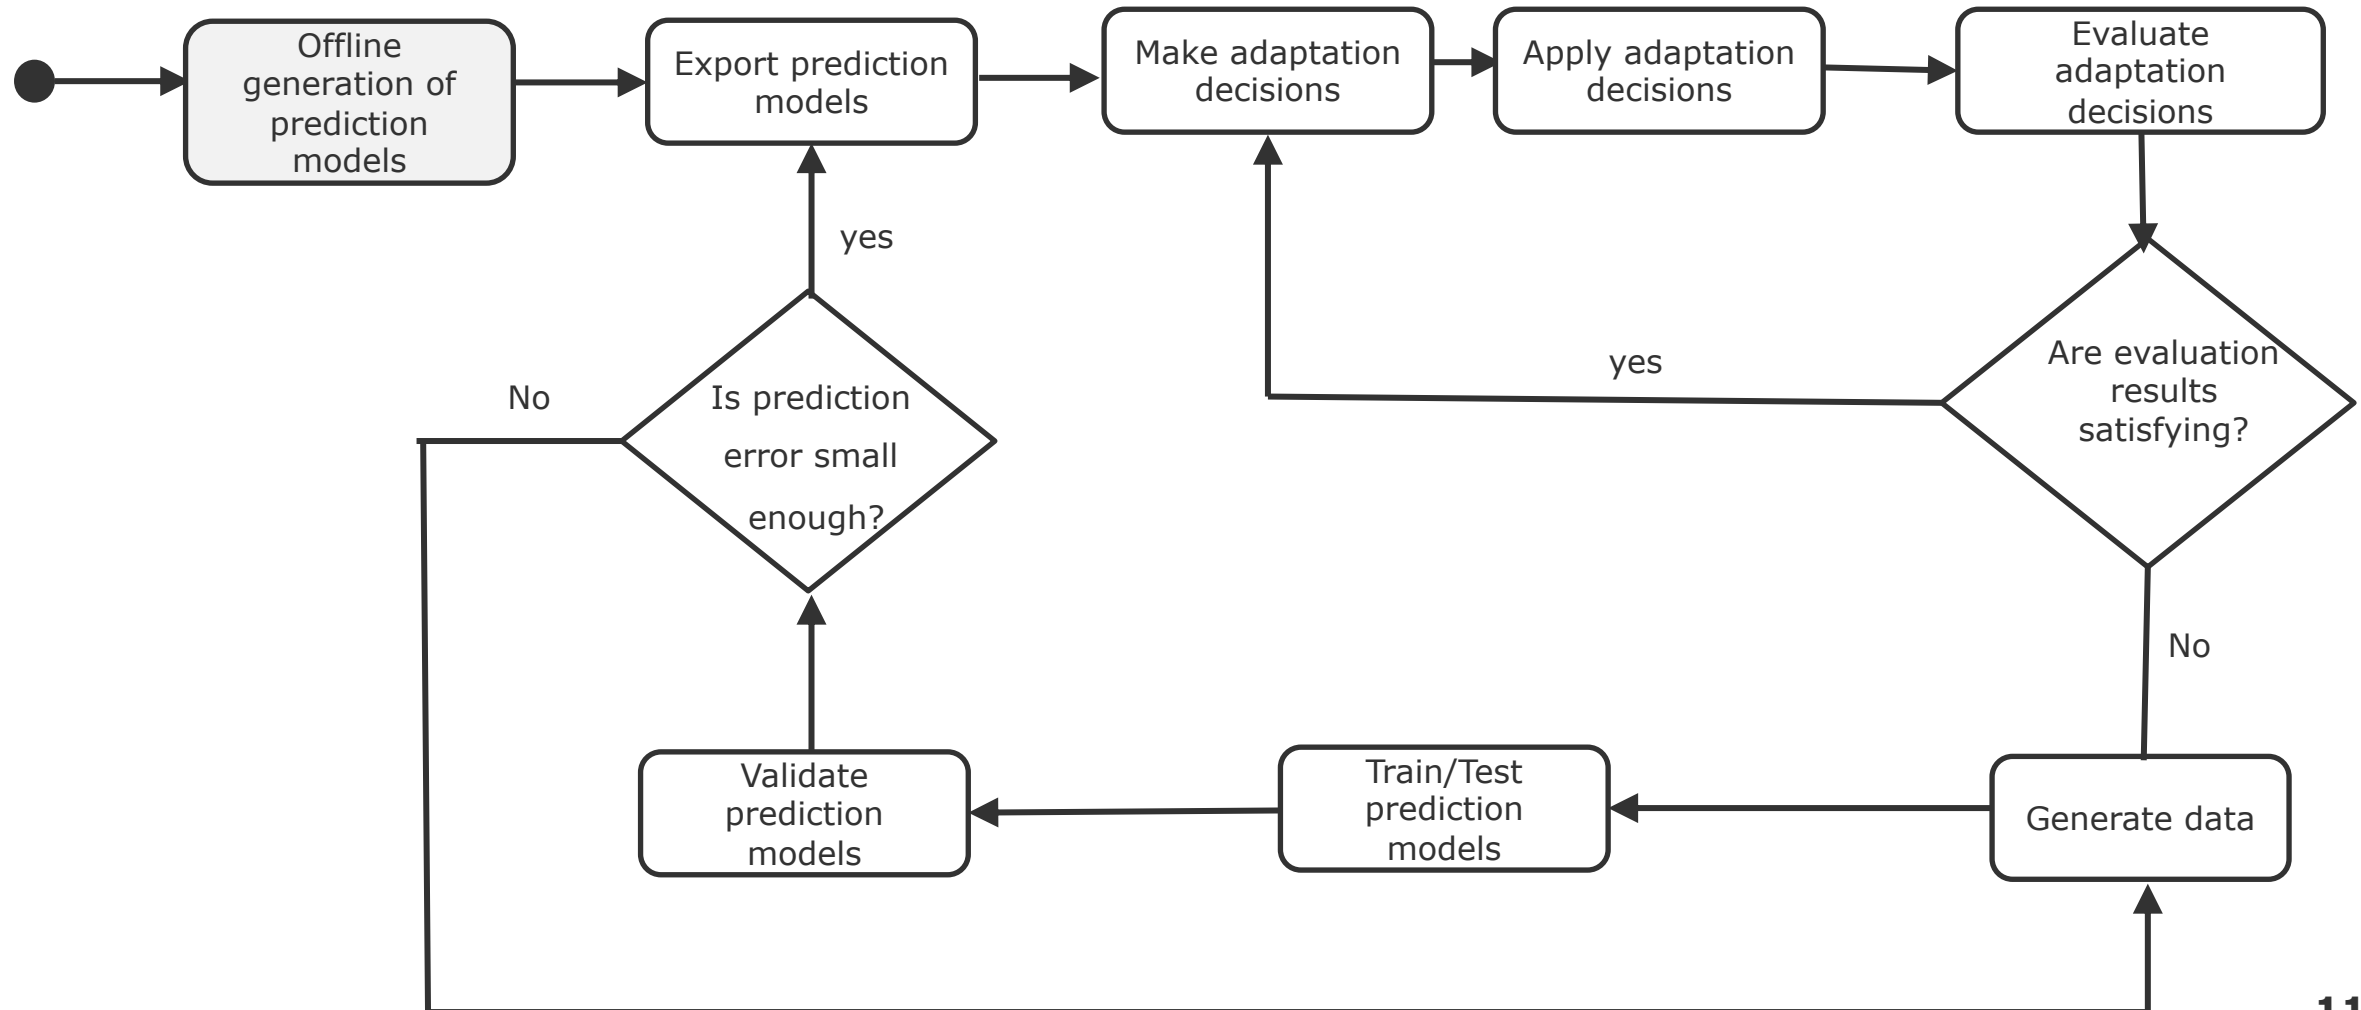
Online Learning for Self-healing and self-optimization



Offline Learning (detailed look)



Online Learning (detailed look)



Overview of Scenario and Models

Elements of Reinforcement Learning

- Exploration & Exploitation - Multi-armed bandits
- Markov Decision Process (MDP)
- Markov property, Bellman Equation, Optimality

**Concepts and
Theory and
Assumptions**

Model-Based methods

- Finite MDP
- Dynamic Programming (Policy and Value Iteration)

**Tabular Methods
Finite/Small State
Spaces**

Model-Free methods

- Monte Carlo
- Temporal Difference (On-policy/Sarsa, Off-policy/Q-Learning)


Approximate methods

- Generalized solution (Incremental, Batch)
- Parameterized solution (Gradient Policy)

**Function
Approximation
Infinite/Large State
Spaces**

Sutton, R. S., & Barto, A. G. (2015). ***Reinforcement learning: An introduction***. MIT press.


free access at the author's page: <http://richsutton.com>



Aleksandrs Slivkins (2019). ***Introduction to Multi-Armed Bandits***, Foundations and Trends® in Machine Learning: Vol. 12: No. 1-2, pp 1-286

free access by the authors'

<https://arxiv.org/abs/1904.07272>



Phase-1 Environment setup with Multi-armed Bandit

Phase-2 Model-Based implementation

Phase-3 Model-Free extension

Phase-4 Approximate model extension

"When solving a problem of interest, do not solve a more general problem as an intermediate step"

Vladimir Vapnik [1]

General goal

For each phase, you will receive:

- a working algorithm that integrates with some simple environment
- a training version of the environment to train your algorithm
- a testing version of the environment to test your algorithm

Your goal will be to:

- Document the algorithm
- Make any needed extensions to integrate with new environment
- Run experiments and plot the results (error approximation by episodes)

Event: One of more components failing


Adaptive response: Select the repair actions (rules) that more quickly restore the system

Challenges:

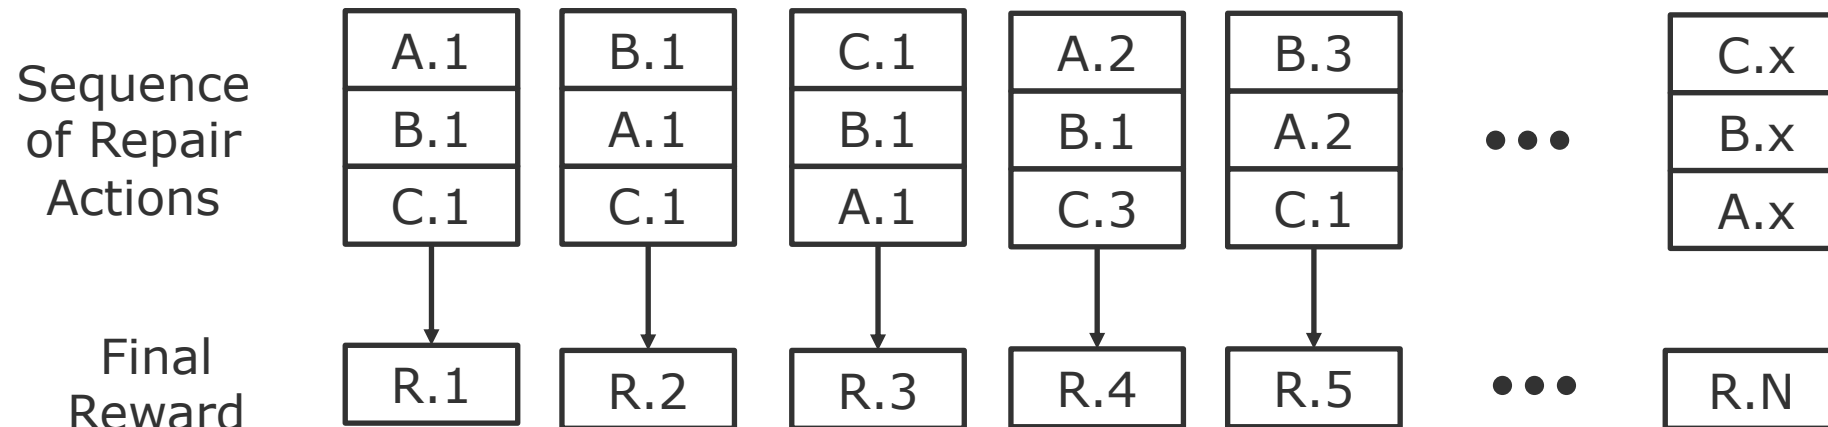
1. Multiple components might fail at the same time
2. For the same failure, there might be multiple actions that repair a component

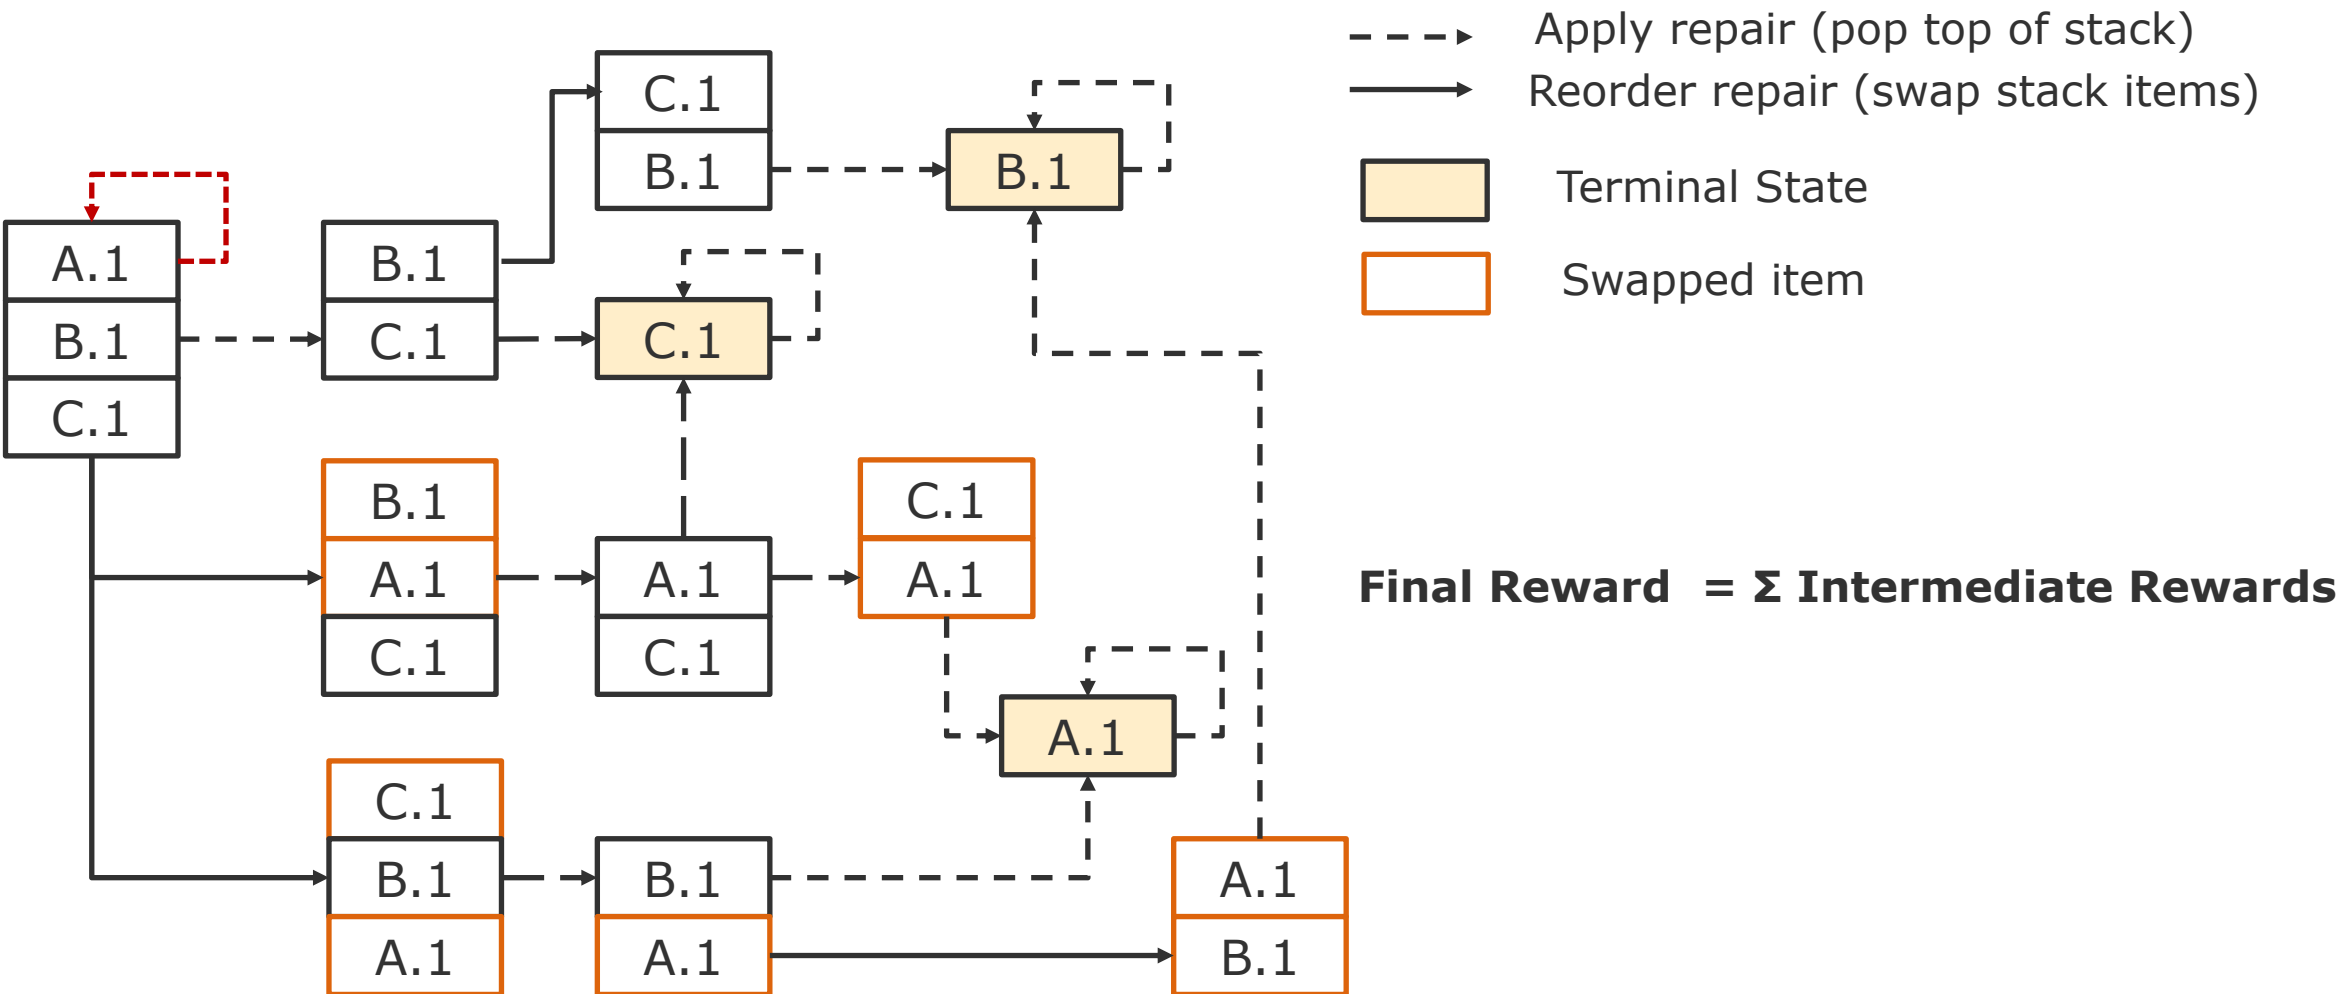
Goal:

Find the best sequence of repair actions



Multi-Armed Bandit Model





End